

## ASSIGNMENT 1: Lexical Analysis

## OVERVIEW

The first major component of your compiler is the lexical analyzer. It is recommended that you build this using the lex/flex tool. You may also code it by hand, **however this is not a productive use of your valuable time.**

You will also need to write a small testing program to try out your lexer. Finally, you will apply your lexer and test program to a set of test cases, and compare the output to "correct" results.

## INPUT ASSUMPTIONS

In the C language, pre-processing takes place before lexical analysis, although many compilers (including gcc) perform both within the same lexer. You may assume that the input to your lexer has already been passed through the preprocessor, which can be invoked with the command gcc -E. Pre-processing handles the following for you:

- include files
- conditional inclusion of source (e.g. #ifdef)
- trigraph conversion and other stupidity (but not C-99 "digraphs")
- joining of lines continued with a \ at the end of the line
- macros

According to section 5.1.1.2 of the ISO C Standard, this corresponds to translation phases 1 through 4. **You don't need to worry about pre-processing, other than picking up the line numbers:**

## LINE NUMBERS AND ERROR REPORTING:

You will see that the output of gcc -E contains lines which begin with a #. These are line number markers, and tell you where in the original source file the next line came from. These markers are output by gcc -E whenever it takes any action which adds or deletes lines. By extracting the information from these marker lines, and keeping track of newlines, you can maintain the current filename and line number for error reporting purposes.

Error reporting and in particular error recovery is not easy, but we're going to get started with this first assignment. If your lexer encounters input which does not conform to the lexical rules, it should output a descriptive error message with the file name and line number of the place of error. It may then halt. As a future refinement, we'll discuss ways attempt recovery and further processing. Remember: error messages go to stderr, not stdout! Example:

```
$ ./mylexer test1.c
```

```
test1.c:10: Error: unrecognized character '@'
```

## TOKEN CODES

yacc/bison expects yylex() to return integer token codes. By convention, single-character tokens can be represented simply by their ASCII value. E.g. if you recognize an = token, you may return('='). This obviously won't work for multi-character tokens. You will need to provide a .h file containing a list of #defines for each token code that is not represented by a simple character. yacc/bison likes these numbers to start at 257 to avoid conflict with ASCII codes, e.g. #define EQEQ 257 #define NOTEQ 258

Note that yacc/bison prefers to see 0 as the EOF token, not -1, but it tolerates -1.

yacc/bison can generate this .h file for you, or you can create it manually and have it included with your parser. For this assignment, we are not using a parser, and so you will have to manually create the file (or, for testing purposes, use the tokens-manual.h which is provided).

## WHAT ARE THE TOKENS?/ WHAT CAN I SKIP?

Read the C standard or Harbison & Steele. Basically, tokens are keywords, identifiers, numeric constants, string literals, or operators/punctuators.

You don't need to implement Universal Character Names.

You must implement all of the numeric constants, including properly recognizing them and converting them to an internal representation (int, float, etc.). See below about token values. You should recognize floating-point constants with a "p" form but aren't required to understand their value (although that isn't at all difficult). You should recognize an L in front of a character or string, but aren't required to understand or deal with wide characters. *If you find that you are pressed for time, it is better to turn in something with broken floating-point (real numbers) handling than to fall behind. We will not be dealing with floating point arithmetic in the final compiler.*

Recognize all of the keywords listed in the standard. If you want to skip "digraphs", which are a C99 extension for programmers unfortunate enough to have broken keyboards, you may, although they aren't difficult either.

Note that whitespace is not a token! Whitespace occurs between tokens, and is silently consumed.

## TOKEN SEMANTIC VALUES

yacc/bison expects the lexer to return token semantic values in a global variable called `yyval`. It doesn't look at this variable to make parse decisions, but as we'll see, being able to maintain semantic values for terminals and non-terminals in the grammar is the underpinning for the entire syntax-directed translation process.

`yyval` is by default an int but that isn't very useful for anything but a trivial calculator. Let's say that lexical values could be either strings or integers. yacc/bison is expecting a union definition like this:

```
typedef union {
    char *string_literal;
    int integer;
} YYSTYPE;
```

```
extern YYSTYPE yyval;
```

The above is representative of the form of the declaration although you will probably need something a bit more complicated. It is up to you to provide this declaration in your header file, which will ultimately be included by your parser. For now, you can return identifiers simply as a `char *`. In the future, you'll see how it is helpful to have a symbol table as part of your lexer and identifiers will be returned as pointers to a struct of your design. For numeric constants, you need to somehow indicate what *type* of constant you have seen. This is discussed below.

**Note:** It is a good idea to pass all information about the lexical semantic value of a token in the single global `yyval` union. While it is tempting to have other global variables for this purpose, and especially tempting in this first assignment, this approach will cause problems later on!

## TESTING

To test your lexer, write a simple main program which repeatedly calls `yylex()` and prints out **one line for each token**, delimited by **tabs**, and containing the following columns:

filename in which the token was found

line number at which the token was found  
 token name  
 token semantic value  
 other token information (if necessary)

As we'll discuss in class, by standardizing on the output format, we will be able to compare and contrast the performance of each student's lexer, and/or to compare the output of your lexer against a presumed good output.

filename: The name of the file at which the token was found. This information can be extracted from the # directives if the input is coming from the preprocessor. If you have not yet seen any such directives, use <stdin>

line number: an integer formatted as %d

token name: for single-character tokens, use the actual character, e.g. [. For multi-character or variable-length tokens, you will have a defined name for each. To aid testing, please use the names found at the end of this document. They should be self-explanatory.

token semantic value: don't output this column (including the \t) if there is no value associated with the token, such as punctuation marks. For all other token types, output according to the following:

| Token Type | Output  |
|------------|---|
| IDENT      | The identifier name   |
| NUMBER     | value, %lld for ints, %Lg for reals   |
| CHARLIT    | The actual character, if it is printable, otherwise the escape code:<br>\0, \a, \b, \f, \n, \r, \t, \v<br>else: \\%03o<br>also, always escape ' " and \           |
| STRING     | Each character of the string, except for the implicitly added nul terminator, all escaped if necessary as with CHARLIT.<br>Ponder: "\0\0\0" is not the same as "" |

For CHARLIT and STRING token types, the token value is the single byte or the array of bytes, respectively. It is **not** the original lexeme. There are many different ways to represent the same character or string literal. They should be reduced to an identical internal representation. Write the code to convert the character literal to a char, or the string to an array of chars terminated by a . This includes all of the escape code processing. Then write the code to display a single character, or a string of characters, including the escape conventions for output described above. Do **not** simply print out yytext to show the semantic value of a STRING or CHARLIT.

Handling all cases of character literals and string literals can be somewhat difficult. Suggestion: read up on "start conditions" in lex/flex. Note that an elegant solution will use the same logic and code to handle character escapes, regardless of whether they are seen inside a character literal or a string literal.

For some token types which are numbers, output a 5th column with additional information about the type of the number as inferred from its lexical form. E.g. a decimal constant followed by the letters ULL tells us that it is an unsigned long long. Although we are a long way from discussing the C type system, it is important that we capture these lexical type cues now. For real NUMBERS, output the size specifier: FLOAT, DOUBLE, LONGDOUBLE.

For integer NUMBERS, output UNSIGNED, if the constant is tagged as unsigned, then output the size specification: INT, LONG, LONGLONG.

As an example, here is part of the output from the test cases:

|                |   |        |                        |     |                    |
|----------------|---|--------|------------------------|-----|--------------------|
| ltests/chars.c | 1 | STRING | \\0\\0\\0              |     |                    |
| ltests/chars.c | 5 | STRING | \\0018\\377\\001\\256Z |     |                    |
| ltests/kw.c    | 1 | AUTO   |                        |     |                    |
| ltests/kw.c    | 1 | ;      |                        |     |                    |
| ltests/num.c   | 5 | NUMBER | REAL                   | 0.1 | DOUBLE             |
| ltests/num.c   | 8 | NUMBER | INTEGER                | 63  | UNSIGNED, LONGLONG |

### TEST CASES

You'll find a few .c files in the subdirectory ltests on the course web site. You can run all of them through your testing program, e.g.

```
gcc -E ltests/*.c | ./lexertester
```

The files ltests.out and ltests.err contain example output and errors from doing so. Note that the test cases contain a few tricky conditions and warnings/errors, as discussed in class.

### TOKEN NAMES

The token names that were used in the test case output are defined in the file tokens-manual.h which is on the course website. They are an enum although you could also use #define. Note that single-character tokens such as + are represented by their ASCII value and are not in this list.